

COMP9313: Big Data Management



Lecturer: Xin Cao

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Chapter 4.1: Spark I

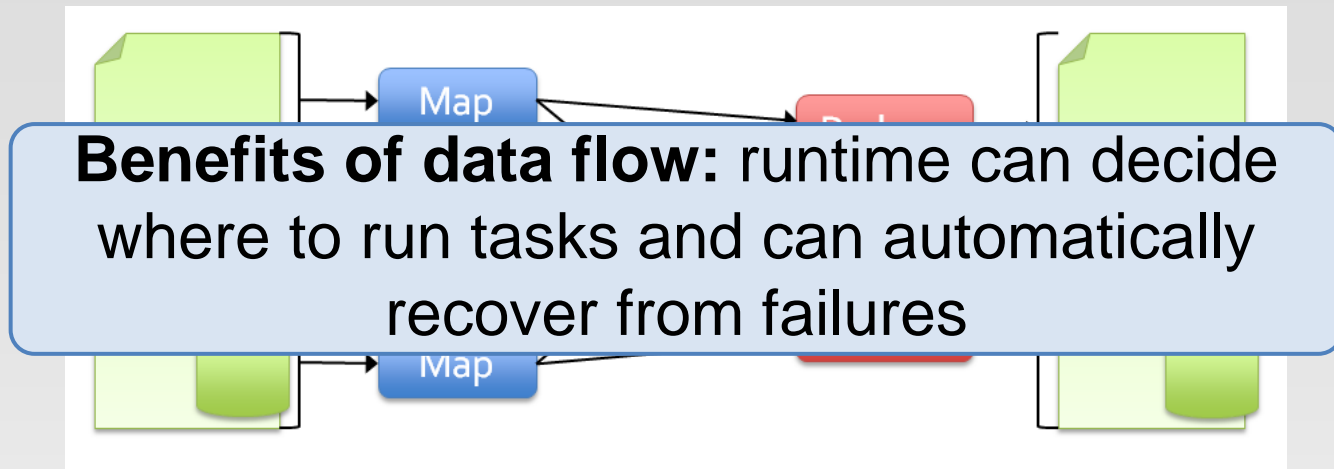


Part 1: Spark Introduction

Limitations of MapReduce

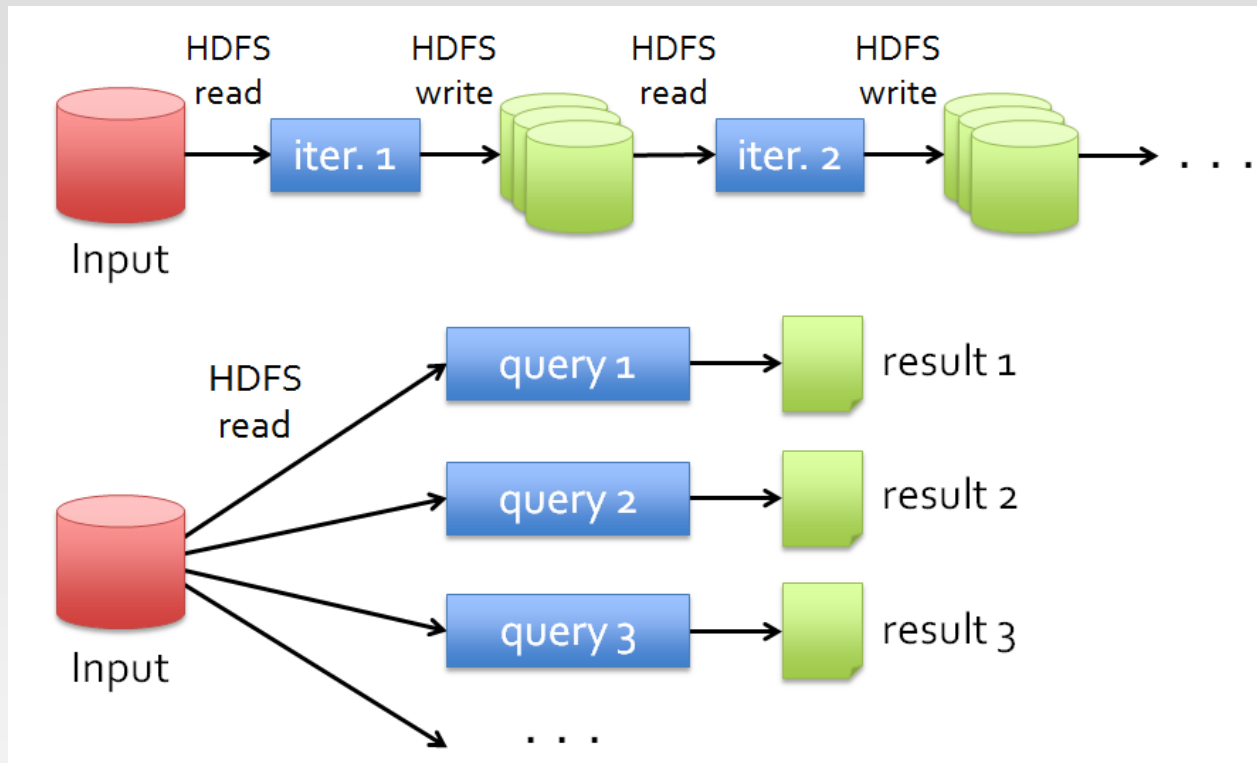
- ❖ MapReduce greatly simplified big data analysis on large, unreliable clusters. It is great at one-pass computation.
- ❖ But as soon as it got popular, users wanted more:
 - More **complex**, multi-pass analytics (e.g. ML, graph)
 - More **interactive** ad-hoc queries
 - More **real-time** stream processing
- ❖ All 3 need faster **data sharing** across parallel jobs
 - One reaction: specialized models for some of these apps, e.g.,
 - ▶ Pregel (graph processing)
 - ▶ Storm (stream processing)

Limitations of MapReduce



- ❖ As a general programming model:
 - It is more suitable for one-pass computation on a large dataset
 - Hard to compose and nest multiple operations
 - No means of expressing iterative operations
- ❖ As implemented in Hadoop
 - All datasets are read from disk, then stored back on to disk
 - All data is (usually) triple-replicated for reliability
 - Not easy to write MapReduce programs using Java

Data Sharing in MapReduce




Slow due to replication, serialization, and disk IO

- ❖ Complex apps, streaming, and interactive queries all need one thing that MapReduce lacks:

Efficient primitives for **data sharing**

Data Sharing in MapReduce

- ❖ Iterative jobs involve a lot of disk I/O for each repetition




- ❖ Interactive queries and online processing involves lots of disk I/O



Hardware for Big Data



Lots of hard drives



Lots of CPUs




And lots of memory!

Goals of Spark

- ❖ Keep more data in-memory to improve the performance!
- ❖ Extend the MapReduce model to better support two common classes of analytics apps:
 - Iterative algorithms (machine learning, graphs)
 - Interactive data mining
- ❖ Enhance programmability:
 - Integrate into Scala programming language
 - Allow interactive use from Scala interpreter


Data Sharing in Spark Using RDD



10-100 × faster than network and disk

What is Spark

- ❖ One popular answer to “What’s beyond MapReduce?”
- ❖ Open-source engine for large-scale distributed data processing
 - Supports generalized dataflows
 - Written in Scala, with bindings in Java, Python, and R
- ❖ Brief history:
 - Developed at UC Berkeley AMPLab in 2009
 - Open-sourced in 2010
 - Became top-level Apache project in February 2014
 - Commercial support provided by DataBricks



What is Spark

- ❖ Fast and expressive cluster computing system interoperable with Apache Hadoop
- ❖ Improves efficiency through:
 - **In-memory** computing primitives
 - General computation graphs→ Up to 100 × faster
(10 × on disk)
- ❖ Improves usability through:
 - Rich APIs in Scala, Java, Python
 - Interactive shell→ Often 5 × less code

What is Spark

❖ Spark is not

- a modified version of Hadoop
- dependent on Hadoop because it has its own cluster management
- Spark uses Hadoop for storage purpose only

❖ Spark's design philosophy centers around four key characteristics:

- Speed
- Ease of use
- Modularity
- Extensibility

Speed

- ❖ Its internal implementation benefits immensely from the performance improvement of CPUs and memory.
 - The framework is optimized to take advantage of memory, multiple cores, and the underlying Unix-based operating system
- ❖ Spark builds its query computations as a directed acyclic graph
 - Tasks can execute in parallel across workers on the cluster
- ❖ It has a physical execution engine which generates compact code for execution

Ease of Use

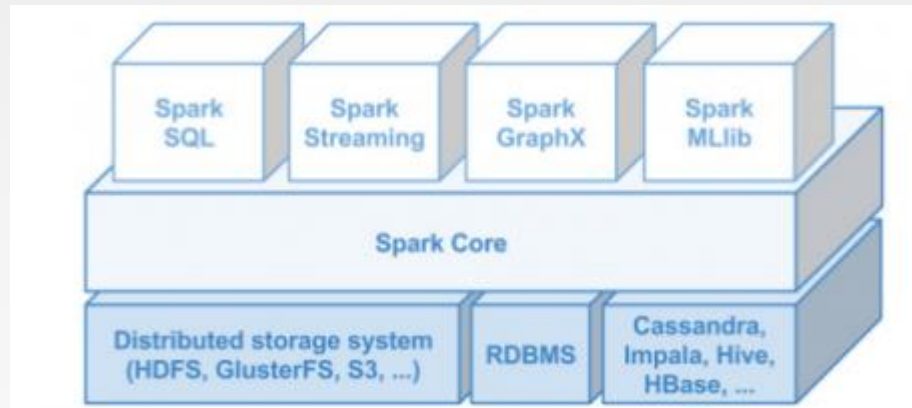
- ❖ Spark achieves simplicity by providing a fundamental abstraction of a simple logical data structure called a Resilient Distributed Dataset (RDD)
- ❖ Since Spark 2.x, DataFrames and Datasets APIs have been developed upon RDD
- ❖ By providing a set of ***transformations*** and ***actions*** as operations, Spark offers a simple programming model that you can use to build big data applications in familiar languages.

Modularity

- ❖ Spark operations can be applied across many types of workloads and expressed in any of the supported programming languages: Scala, Java, Python, SQL, and R.
- ❖ Spark offers unified libraries with well-documented APIs that include the following modules as core components: Spark SQL, Spark Structured Streaming, Spark MLlib, and GraphX, combining all the workloads running under one engine.
- ❖ You can write a single Spark application that can do it all—no need for distinct engines for disparate workloads, no need to learn separate APIs.

Extensibility

- ❖ Spark focuses on its fast, parallel computation engine rather than on storage.
 - You can use Spark to read data stored in myriad sources—local file systems, Apache Hadoop, Apache Cassandra, Apache HBase, MongoDB, Apache Hive, RDBMSs, and more—and process it all in memory.
- ❖ Spark's DataFrameReaders and DataFrameWriters can also be extended to read data from other sources, such as Apache Kafka, Kinesis, Azure Storage, and Amazon S3



What is Spark

- ❖ Spark is the basis of a wide set of projects in the Berkeley Data Analytics Stack (BDAS)

Spark SQL
(SQL)

Spark
Streaming
(real-time)

GraphX
(graph)

MLlib
(machine
learning)


...

Spark Core
(Scala, Python, Java, R, SQL)

- Spark SQL (SQL on Spark)
- Spark Streaming (stream processing)
- GraphX (graph processing)
- MLlib (machine learning library)

Spark's Ecosystem of Connectors


- ❖ The community of Spark developers maintains a list of third-party Spark packages as part of the growing ecosystem



Spark Ideas

- ❖ Expressive computing system, not limited to map-reduce model
- ❖ Facilitate system memory
 - avoid saving intermediate results to disk
 - cache data for repetitive queries (e.g. for machine learning)
- ❖ Layer an in-memory system on top of Hadoop.
- ❖ Achieve fault-tolerance by re-execution instead of replication

Spark Workflow



- ❖ A Spark program first creates a SparkContext object
 - Tells Spark how and where to access a cluster
 - Define RDDs
 - Connect to several types of cluster managers (e.g., YARN, Mesos, or its own manager)
- ❖ Cluster manager:
 - Allocate resources across applications
- ❖ Spark executor:
 - Run computations
 - Access data storage

Download and Configure Spark

- ❖ Current version: 3.3.0. <https://spark.apache.org/downloads.html>
 - You also need to install Java first

Download Apache Spark™

1. Choose a Spark release: 3.3.0 (Jun 16 2022) ▼
2. Choose a package type: Pre-built for Apache Hadoop 3.3 and later ▼
3. Download Spark: [spark-3.3.0-bin-hadoop3.tgz](#)
4. Verify this release using the 3.3.0 [signatures](#), [checksums](#) and [project release KEYS](#) by following these [procedures](#).

Note that Spark 3 is pre-built with Scala 2.12 in general and Spark 3.2+ provides additional pre-built distribution with Scala 2.13.

- ❖ After downloading the package, unpack it and then configure the path variable in file ~/.bashrc

```
export SPARK_HOME=/home/comp9313/spark
export PATH=$SPARK_HOME/bin:$PATH
```

Spark Shell


- ❖ Spark comes with four widely used interpreters that act like interactive “shells” and enable ad hoc data analysis: pyspark, spark-shell, sparksql, and sparkR

[illegible]

Word Count in Spark (Python)

```
textfile = sc.textFile("hdfs://...", 4)

words = textfile.flatMap(lambda line: line.split(" "))
pairs = words.map(lambda word: (word, 1))
count = pairs.reduceByKey(lambda a, b: a + b)
count.collect()
```

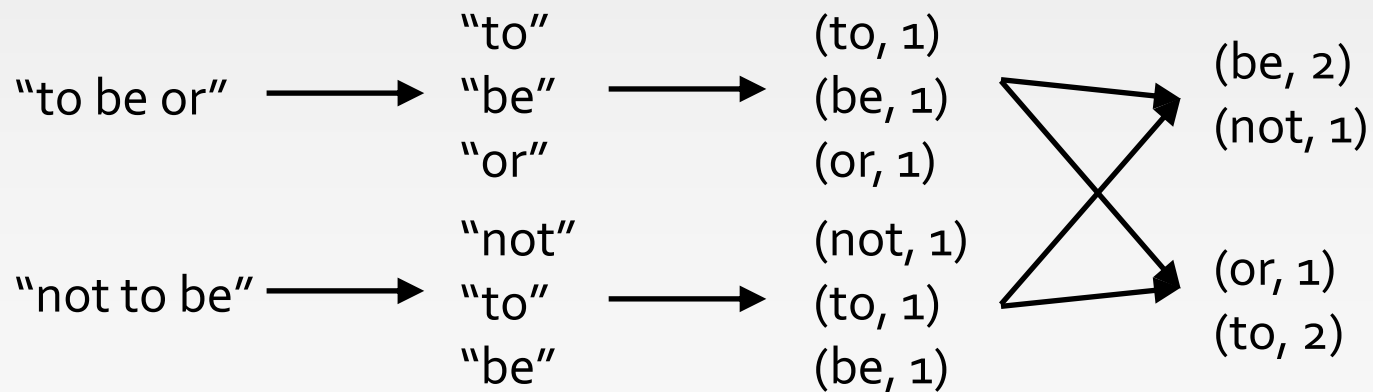


Word Count in Spark (Scala)

```
val file = sc.textFile("hdfs://...")

val counts = file.flatMap(line => line.split(" "))
                  .map(word => (word, 1))
                  .reduceByKey(_ + _)

counts.saveAsTextFile("hdfs://...")
```



Part 2: Scala Introduction

Scala (Scalable language)

- ❖ Scala is a *general-purpose programming language* designed to express common programming patterns in a concise, elegant, and type-safe way
- ❖ Scala supports both Object Oriented Programming and Functional Programming
- ❖ Scala is Practical
 - Can be used as drop-in replacement for Java
 - ▶ Mixed Scala/Java projects
 - Use existing Java libraries
 - Use existing Java tools (Ant, Maven, JUnit, etc...)
 - Decent IDE Support (NetBeans, IntelliJ, Eclipse)



Why Scala

- ❖ Scala supports object-oriented programming. Conceptually, every value is an object and every operation is a method-call. The language supports advanced component architectures through classes and traits
- ❖ Scala is also a functional language. Supports functions, immutable data structures and preference for immutability over mutation
- ❖ Seamlessly integrated with Java
- ❖ Being used heavily for Big data, e.g., Spark, Kafka, etc.

Scala Basic Syntax

- ❖ When considering a Scala program, it can be defined as a collection of objects that communicate via invoking each other's methods.
- ❖ **Object** – same as in Java
- ❖ **Class** – same as in Java
- ❖ **Methods** – same as in Java
- ❖ **Fields** – Each object has its unique set of instant variables, which are called fields. An object's state is created by the values assigned to these fields.
- ❖ **Traits** – Like Java Interface. A trait encapsulates method and field definitions, which can then be reused by mixing them into classes.
- ❖ **Closure** – A **closure** is a function, whose return value depends on the value of one or more variables declared outside this function.

closure = function + environment

Object-Oriented Programming in Scala

- ❖ Scala is object-oriented, and is based on Java's model
- ❖ An **object** is a **singleton object** (there is only one of it)
 - Variables and methods in an **object** are somewhat similar to Java's **static** variables and methods
 - Reference to an **object**'s variables and methods have the syntax ***ObjectName.methodOrVariableName***
 - The name of an **object** should be capitalized
- ❖ A **class** may take parameters, and may describe any number of objects
 - The class body *is* the constructor, but you can have additional constructors
 - With correct use of **val** and **var**, Scala provides getters and setters for class parameters

Scala is Statically Typed

- ❖ You don't have to specify a type in most cases
- ❖ Type Inference

```
val sum = 1 + 2 + 3
```

```
val nums = List(1, 2, 3)
```

```
val map = Map("abc" -> List(1,2,3))
```

Explicit Types

```
val sum: Int = 1 + 2 + 3
```

```
val nums: List[Int] = List(1, 2, 3)
```

```
val map: Map[String, List[Int]] = ...
```

Scala is High level

// Java - Check if string has uppercase character

```
boolean hasUpperCase = false;
for(int i = 0; i < name.length(); i++) {
    if(Character.isUpperCase(name.charAt(i))) {
        hasUpperCase = true;
        break;
    }
}
```

// Scala

```
val hasUpperCase = name.exists(_.isUpper)
```


Scala is Concise

// Java

```
public class Person {  
    private String name;  
    private int age;  
    public Person(String name, Int age) {  
        this.name = name;  
        this.age = age;  
    }  
    public String getName() {                // name getter  
        return name;  
    }  
    public int getAge() {                    // age getter  
        return age;  
    }  
    public void setName(String name) {      // name setter  
        this.name = name;  
    }  
    public void setAge(int age) {           // age setter  
        this.age = age;  
    }  
}
```

// Scala

```
class Person(var name: String, private var _age: Int) {  
    def age = _age                // Getter for age  
    def age_=(newAge: Int) {      // Setter for age  
        println("Changing age to: "+newAge)  
        _age = newAge  
    }  
}
```

Variables and Values

❖ Variables: values stored can be changed

```
var foo = "foo"
```

```
foo = "bar" // okay
```

❖ Values: immutable variable

```
val foo = "foo"
```

```
foo = "bar" // nope
```

Scala is Pure Object Oriented

// Every value is an object

1.toString

// Every operation is a method call

1 + 2 + 3 → (1).+(2).+(3)

// Can omit . and ()

"abc" charAt 1 → "abc".charAt(1)

// Classes (and abstract classes) like Java

```
abstract class Language(val name:String) {  
    override def toString = name  
}
```

// Example implementations

```
class Scala extends Language("Scala")
```

// Anonymous class

```
val scala = new Language("Scala") { /* empty */ }
```

Scala Traits

```
// Like interfaces in Java
trait JVM {
  // But allow implementation
  override def toString = super.toString+" runs on JVM" }
trait Static {
  override def toString = super.toString+" is Static" }

// Traits are stackable
class Scala extends Language with JVM with Static {
  val name = "Scala"
}

println(new Scala) → "Scala runs on JVM is Static"
```

Scala is Functional

- ❖ First-Class Functions. Functions are treated like objects:
 - passing functions as arguments to other functions
 - returning functions as the values from other functions
 - assigning functions to variables or storing them in data structures

```
// Lightweight anonymous functions
```

```
(x:Int) => x + 1
```

```
// Calling the anonymous function
```

```
val plusOne = (x:Int) => x + 1
```

```
plusOne(5) → 6
```

Scala is Functional

- ❖ Closures: a function whose return value depends on the value of one or more variables declared outside this function.

// plusFoo can reference any **values/variables** in scope

```
var foo = 1
```

```
val plusFoo = (x:Int) => x + foo
```

```
plusFoo(5)  →  6
```

// Changing foo changes the return value of plusFoo

```
foo = 5
```

```
plusFoo(5)  →  10
```

Scala is Functional

❖ Higher Order Functions

- A function that does at least one of the following:
 - ▶ takes one or more functions as arguments
 - ▶ returns a function as its result

```
val plusOne = (x:Int) => x + 1
```

```
val nums = List(1,2,3)
```

```
// map takes a function: Int => T
```

```
nums.map(plusOne)      → List(2,3,4)
```

```
// Inline Anonymous
```

```
nums.map(x => x + 1)   → List(2,3,4)
```

```
// Short form
```

```
nums.map(_ + 1)        → List(2,3,4)
```

More Examples on Higher Order Functions

```
val nums = List(1,2,3,4)
```

```
// A few more examples for List class
```

```
nums.exists(_ == 2)      → true
```

```
nums.find(_ == 2)        → Some(2)
```

```
nums.indexOf(_ == 2)     → 1
```

```
// functions as parameters, apply f to the value "1"
```

```
def call(f: Int => Int) = f(1)
```

```
call(plusOne)           → 2
```

```
call(x => x + 1)        → 2
```

```
call(_ + 1)             → 2
```


More Examples on Higher Order Functions

```
val basefunc = (x:Int) => ((y:Int) => x + y)
```

```
// interpreted by:
```

```
basefunc(x){  
    sumfunc(y){ return x+y;}  
    return sumfunc;  
}
```

```
val closure1 = basefunc(1)    closure1(5) = ?  
6
```

```
val closure2 = basefunc(4)    closure2(5) = ?  
9
```

- ❖ basefunc returns a function, and closure1 and closure2 are of function type.
- ❖ While closure1 and closure2 refer to the same function basefunc, the associated environments differ, and the results are different

The Usage of “_” in Scala

- ❖ In anonymous functions, the “_” acts as a placeholder for parameters

```
nums.map(x => x + 1)
```

is equivalent to:

```
nums.map(_ + 1)
```

```
List(1,2,3,4,5).foreach(print(_))
```

is equivalent to:

```
List(1,2,3,4,5).foreach(a => print(a) )
```

- ❖ You can use two or more underscores to refer different parameters.

```
val sum = List(1,2,3,4,5).reduceLeft(_+_)
```

is equivalent to:

```
val sum = List(1,2,3,4,5).reduceLeft((a, b) => a + b)
```

- The reduceLeft method works by applying the function/operation you give it, and applying it to successive elements in the collection

Part 3: RDD Introduction

RDD: Resilient Distributed Datasets

- ❖ Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing. Matei Zaharia, et al. NSDI'12
 - RDD is a **distributed** memory abstraction that lets programmers perform **in-memory** computations on large clusters in a **fault-tolerant** manner.
- ❖ **Resilient**
 - Fault-tolerant, is able to recompute missing or damaged partitions due to node failures.
- ❖ **Distributed**
 - Data residing on multiple nodes in a cluster.
- ❖ **Dataset**
 - A collection of partitioned elements, e.g. tuples or other objects (that represent records of the data you work with).
- ❖ RDD is the primary data abstraction in Apache Spark and the core of Spark. It enables operations on collection of elements in parallel.

RDD: Resilient Distributed Datasets

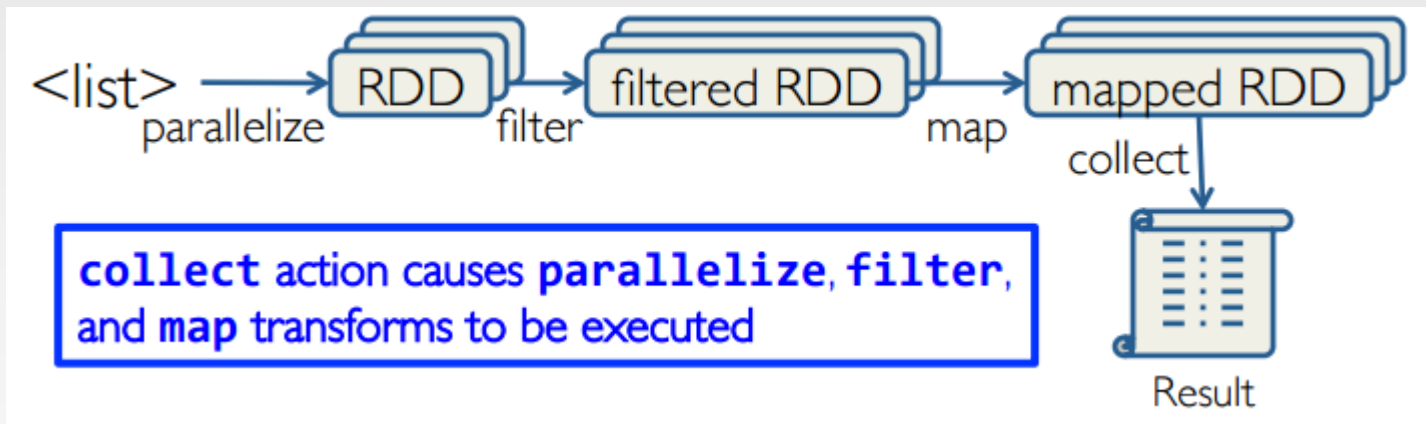
- ❖ *Resilient Distributed Datasets (RDDs)*
 - Distributed collections of objects that can be cached in memory across cluster
 - Manipulated through parallel operators
 - Automatically recomputed on failure based on lineage
- ❖ RDDs can express many parallel algorithms, and capture many current programming models
 - Data flow models: MapReduce, SQL, ...
 - Specialized models for iterative apps: Pregel, ...

RDD Traits

- ❖ **In-Memory**, i.e. data inside RDD is stored in memory as much (size) and long (time) as possible.
- ❖ **Immutable** or **Read-Only**, i.e. it does not change once created and can only be transformed using transformations to new RDDs.
- ❖ **Lazy evaluated**, i.e. the data inside RDD is not available or transformed until an action is executed that triggers the execution.
- ❖ **Cacheable**, i.e. you can hold all the data in a persistent "storage" like memory (default and the most preferred) or disk (the least preferred due to access speed).
- ❖ **Parallel**, i.e. process data in parallel.
- ❖ **Typed**, i.e. values in a RDD have types, e.g. `RDD[Long]` or `RDD[(Int, String)]`.
- ❖ **Partitioned**, i.e. the data inside a RDD is partitioned (split into partitions) and then distributed across nodes in a cluster (one partition per JVM that may or may not correspond to a single node).

Working with RDDs

- ❖ Create an RDD from a data source
 - by parallelizing existing collections (lists or arrays)
 - by transforming an existing RDDs
 - from files in HDFS or any other storage system
- ❖ Apply transformations to an RDD: e.g., map, filter
- ❖ Apply actions to an RDD: e.g., collect, count



- ❖ Users can control two other aspects:
 - Persistence
 - Partitioning

Creating RDDs

- ❖ From HDFS, text files, Amazon S3, Apache HBase, SequenceFiles, any other Hadoop InputFormat

- ❖ Creating an RDD from a File

- `val inputfile = sc.textFile("...", 4)`
 - ▶ RDD distributed in 4 partitions
 - ▶ Elements are lines of input
 - ▶ Lazy evaluation means no execution happens now

```
scala> val inputfile = sc.textFile("pg100.txt")
inputfile: org.apache.spark.rdd.RDD[String] = pg100.txt MapPartitionsRDD[17] at
textFile at <console>:24
```


- ❖ Turn a collection into an RDD

- `sc.parallelize([1, 2, 3])`, creating from a Python list
- `sc.parallelize(Array("hello", "spark"))`, creating from a Scala Array

- ❖ Creating an RDD from an existing Hadoop InputFormat

- `sc.hadoopFile(keyClass, valClass, inputFmt, conf)`

RDD Operations



- ❖ **Transformation:** returns a new RDD.
 - Nothing gets evaluated when you call a Transformation function, it just takes an RDD and return a new RDD.
 - Transformation functions include *map*, *filter*, *flatMap*, *groupByKey*, *reduceByKey*, *aggregateByKey*, *join*, etc.
- ❖ **Action:** evaluates and returns a new value.
 - When an Action function is called on a RDD object, all the data processing queries are computed at that time and the result value is returned.
 - Action operations include *reduce*, *collect*, *count*, *first*, *take*, *countByKey*, *foreach*, *saveAsTextFile*, etc.

Spark Transformations

- ❖ Create new datasets from an existing one
- ❖ Use lazy evaluation: results not computed right away – instead Spark remembers set of transformations applied to base dataset
 - Spark optimizes the required calculations
 - Spark recovers from failures
- ❖ Some transformation functions

Transformation	Description
<code>map(func)</code>	return a new distributed dataset formed by passing each element of the source through a function <i>func</i>
<code>filter(func)</code>	return a new dataset formed by selecting those elements of the source on which <i>func</i> returns true
<code>distinct([numTasks])</code>	return a new dataset that contains the distinct elements of the source dataset
<code>flatMap(func)</code>	similar to map, but each input item can be mapped to 0 or more output items (so <i>func</i> should return a Seq rather than a single item)

Spark Actions

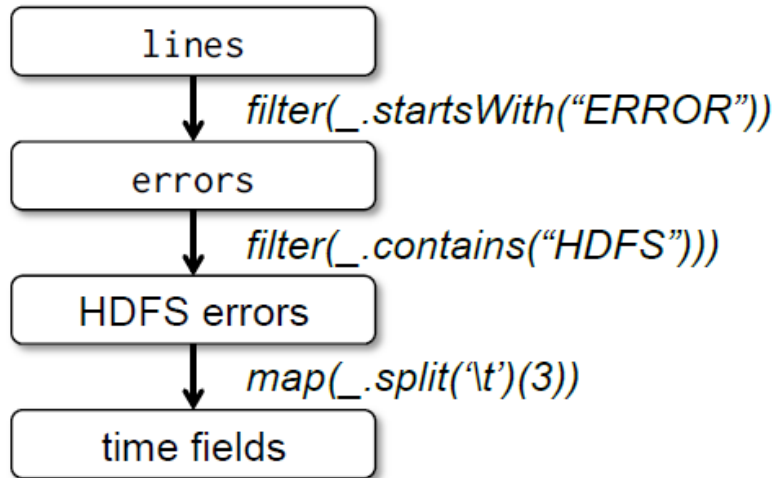
- ❖ Cause Spark to execute recipe to transform source
- ❖ Mechanism for getting results out of Spark
- ❖ Some action functions

Action	Description
<code>reduce(func)</code>	aggregate dataset's elements using function <i>func</i> . <i>func</i> takes two arguments and returns one, and is commutative and associative so that it can be computed correctly in parallel
<code>take(n)</code>	return an array with the first <i>n</i> elements
<code>collect()</code>	return all the elements as an array WARNING: make sure will fit in driver program
<code>takeOrdered(n, key=func)</code>	return <i>n</i> elements ordered in ascending order or as specified by the optional key function

- ❖ Example: `words.collect().foreach(println)`

Example (Scala)

- ❖ Web service is experiencing errors and an operators want to search terabytes of logs in the Hadoop file system to find the cause.



//base RDD

```
val lines = sc.textFile("hdfs://...")
```

//Transformed RDD

```
val errors = lines.filter(_.startsWith("Error"))
```

```
errors.persist()
```

```
errors.count()
```

```
errors.filter(_.contains("HDFS"))
```

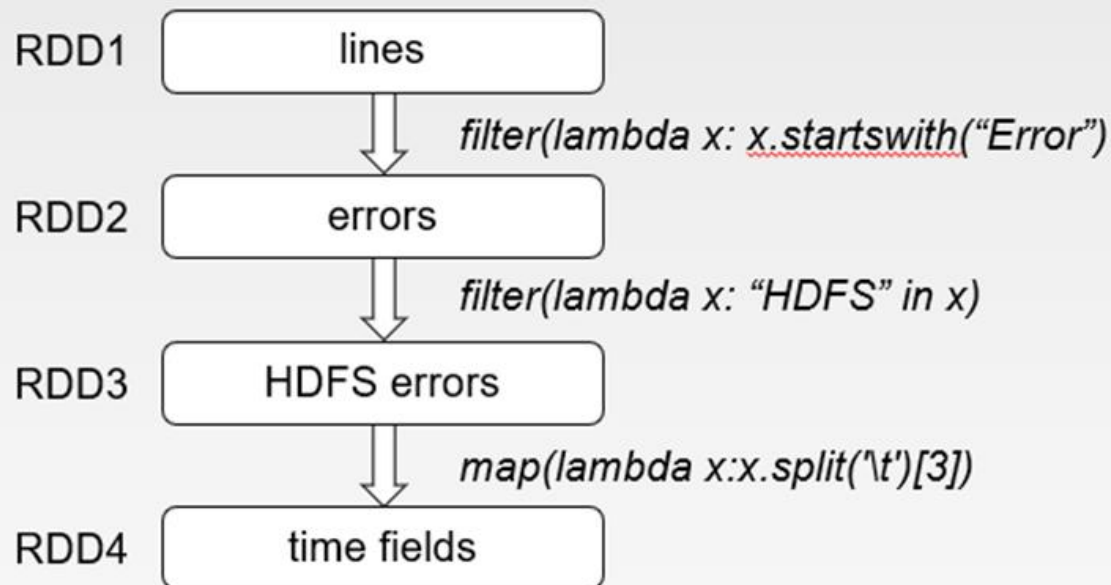
```
.map(_.split('\t')(3))
```

```
.collect()
```

- **Line1:** RDD backed by an HDFS file (base RDD lines not loaded in memory)
- **Line3:** Asks for errors to persist in memory (errors are in RAM)

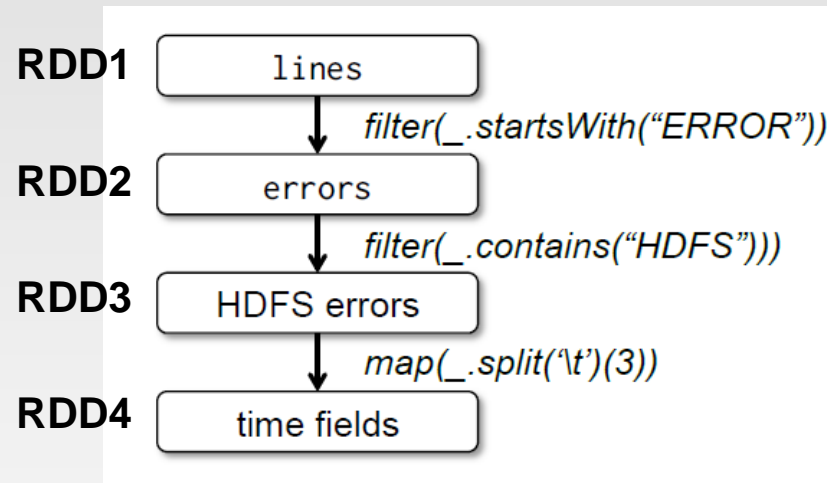
Example (Python)

```
lines = sc.textFile("hdfs://...") //base RDD, obtained from a file on HDFS
errors = lines.filter(lambda x: x.startswith("Error")) //get messages that start
errors.persist() //persist the data in memory
errors.count()
errors.filter(lambda x: "HDFS" in x).map(lambda x:x.split('\t')[3]).collect()
```




Lineage Graph

- ❖ RDDs keep track of lineage
- ❖ RDD has enough information about how it was derived from to compute its partitions from data in stable storage.



- ❖ Example:
 - If a partition of errors is lost, Spark rebuilds it by applying a filter on only the corresponding partition of lines.
 - Partitions can be recomputed in parallel on different nodes, without having to roll back the whole program.

Deconstructed



//base RDD

```
val lines = sc.textFile("hdfs://...")
```

//Transformed RDD

```
val errors = lines.filter(_.startsWith("Error"))
```

```
errors.persist()
```

```
errors.count()
```

```
errors.filter(_.contains("HDFS"))
```

```
.map(_.split('\t')(3))
```


```
.collect()
```

Deconstructed



//base RDD


```
val lines = sc.textFile("hdfs://...")
```



//Transformed RDD

```
val errors = lines.filter(_.startsWith("Error"))
```

```
errors.persist()
```



errors.count()

count() causes Spark to: 1) read data; 2) sum within partitions; 3) combine sums in driver

Put transform and action together:

```
errors.filter(_.contains("HDFS")).map(_split('t')(3)).collect()
```


RDD Persistence: Cache/Persist

- ❖ One of the most important capabilities in Spark is *persisting* (or *caching*) a dataset in memory across operations.
- ❖ When you persist an RDD, each node stores any partitions of it. You can reuse it in other actions on that dataset
- ❖ Each persisted RDD can be stored using a different *storage level*, e.g.
 - MEMORY_ONLY:
 - ▶ Store RDD as deserialized Java objects in the JVM.
 - ▶ If the RDD does not fit in memory, some partitions will not be cached and will be recomputed when they're needed.
 - ▶ This is the default level.
 - MEMORY_AND_DISK:
 - ▶ If the RDD does not fit in memory, store the partitions that don't fit on disk, and read them from there when they're needed.
- ❖ `cache()` = `persist(StorageLevel.MEMORY_ONLY)`

Why Persisting RDD?

```
val lines = sc.textFile("hdfs://...")
```

```
val errors = lines.filter(_.startsWith("Error"))
```

```
errors.persist()
```

```
errors.count()
```

- ❖ If you do `errors.count()` again, the file will be loaded again and computed again.
- ❖ `Persist` will tell Spark to cache the data in memory, to reduce the data loading cost for further actions on the same data
- ❖ `errors.persist()` will do nothing. It is a lazy operation. But now the RDD says "read this file and then cache the contents". The action will trigger computation and data caching.

References

- ❖ <http://spark.apache.org/docs/latest/index.html>
- ❖ <http://www.scala-lang.org/documentation/>
- ❖ <http://www.scala-lang.org/docu/files/ScalaByExample.pdf>
- ❖ [A Brief Intro to Scala](#), by Tim Underwood.
- ❖ [Learning Spark](#). 1st and 2nd Edition

End of Chapter 4.1