


# **COMP9313: Big Data Management**

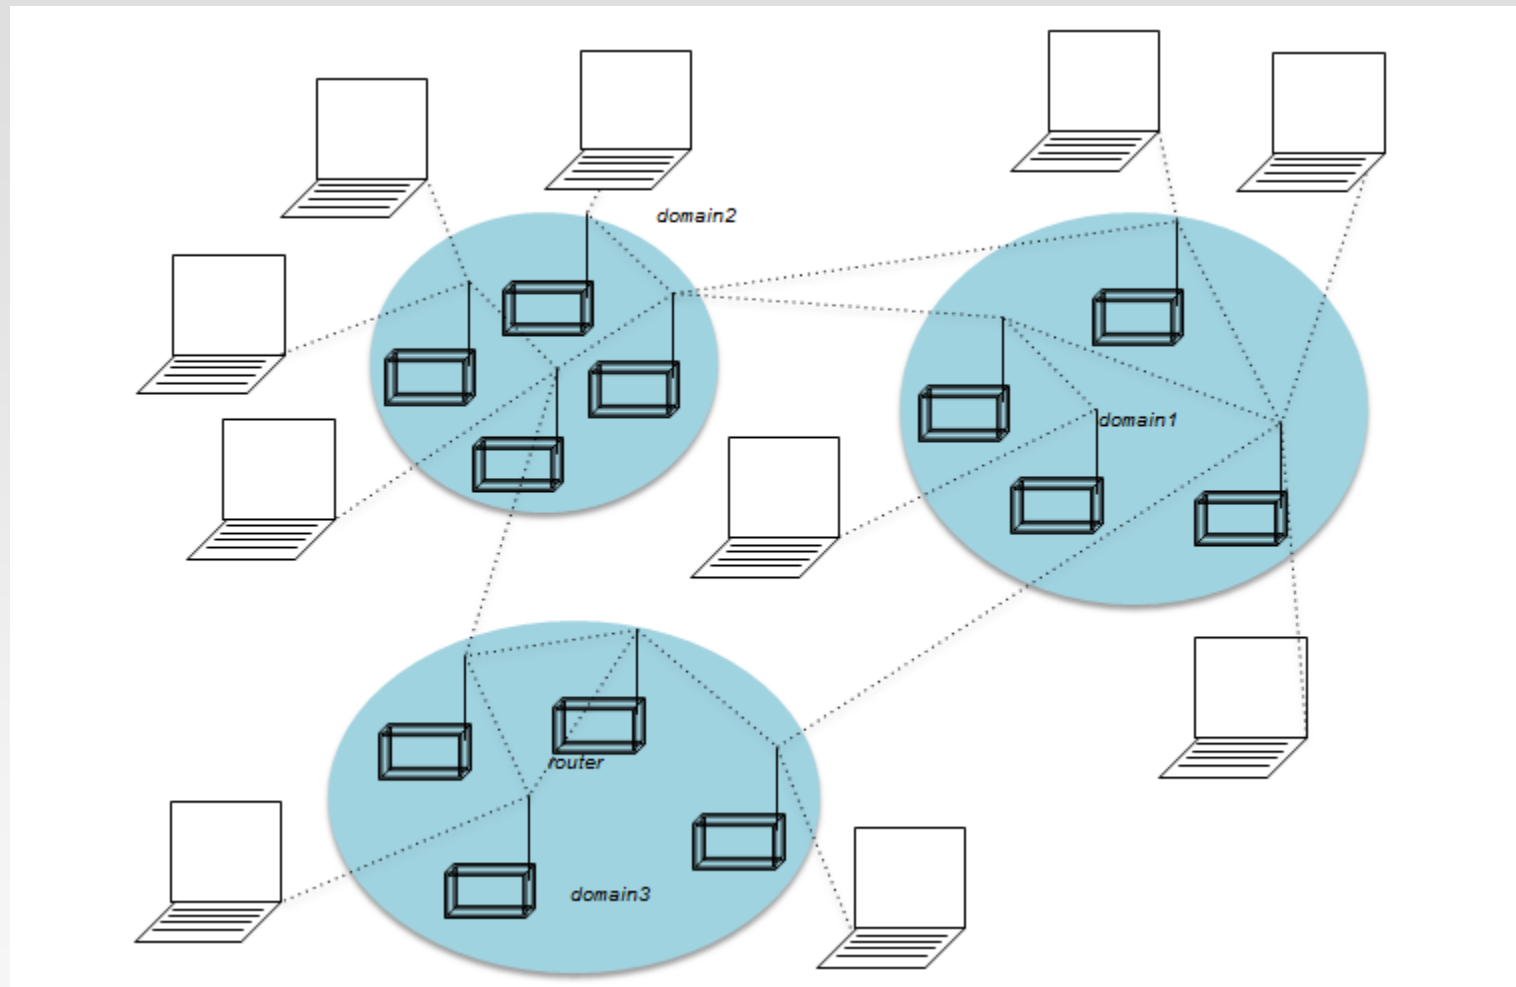


**Lecturer: Xin Cao**

**Course web site: <http://www.cse.unsw.edu.au/~cs9313/>**

# **Chapter 9: Link Analysis - PageRank**


# Graph Data: Communication Nets



**Internet**

# Web as a Directed Graph

- ❖ Web as a directed graph:
  - Nodes: Webpages
  - Edges: Hyperlinks




# Broad Question

- ❖ How to organize the Web?
- ❖ First try: Human curated Web directories
  - Yahoo, LookSmart, etc.
- ❖ Second try: Web Search
  - Information Retrieval investigates:  
Find relevant docs in a small and trusted set
    - ▶ Newspaper articles, Patents, etc.
  - But: Web is huge, full of untrusted documents, random things, web spam, etc.
- ❖ What is the “best” answer to query “newspaper”?
- No single right answer



# Ranking Nodes on the Graph

- ❖ All web pages are not equally “important”
  - <http://xxx.github.io/> vs. <http://www.unsw.edu.au/>
- ❖ There is large diversity in the web-graph node connectivity. Let's rank the pages by the link structure!



# Link Analysis Algorithms

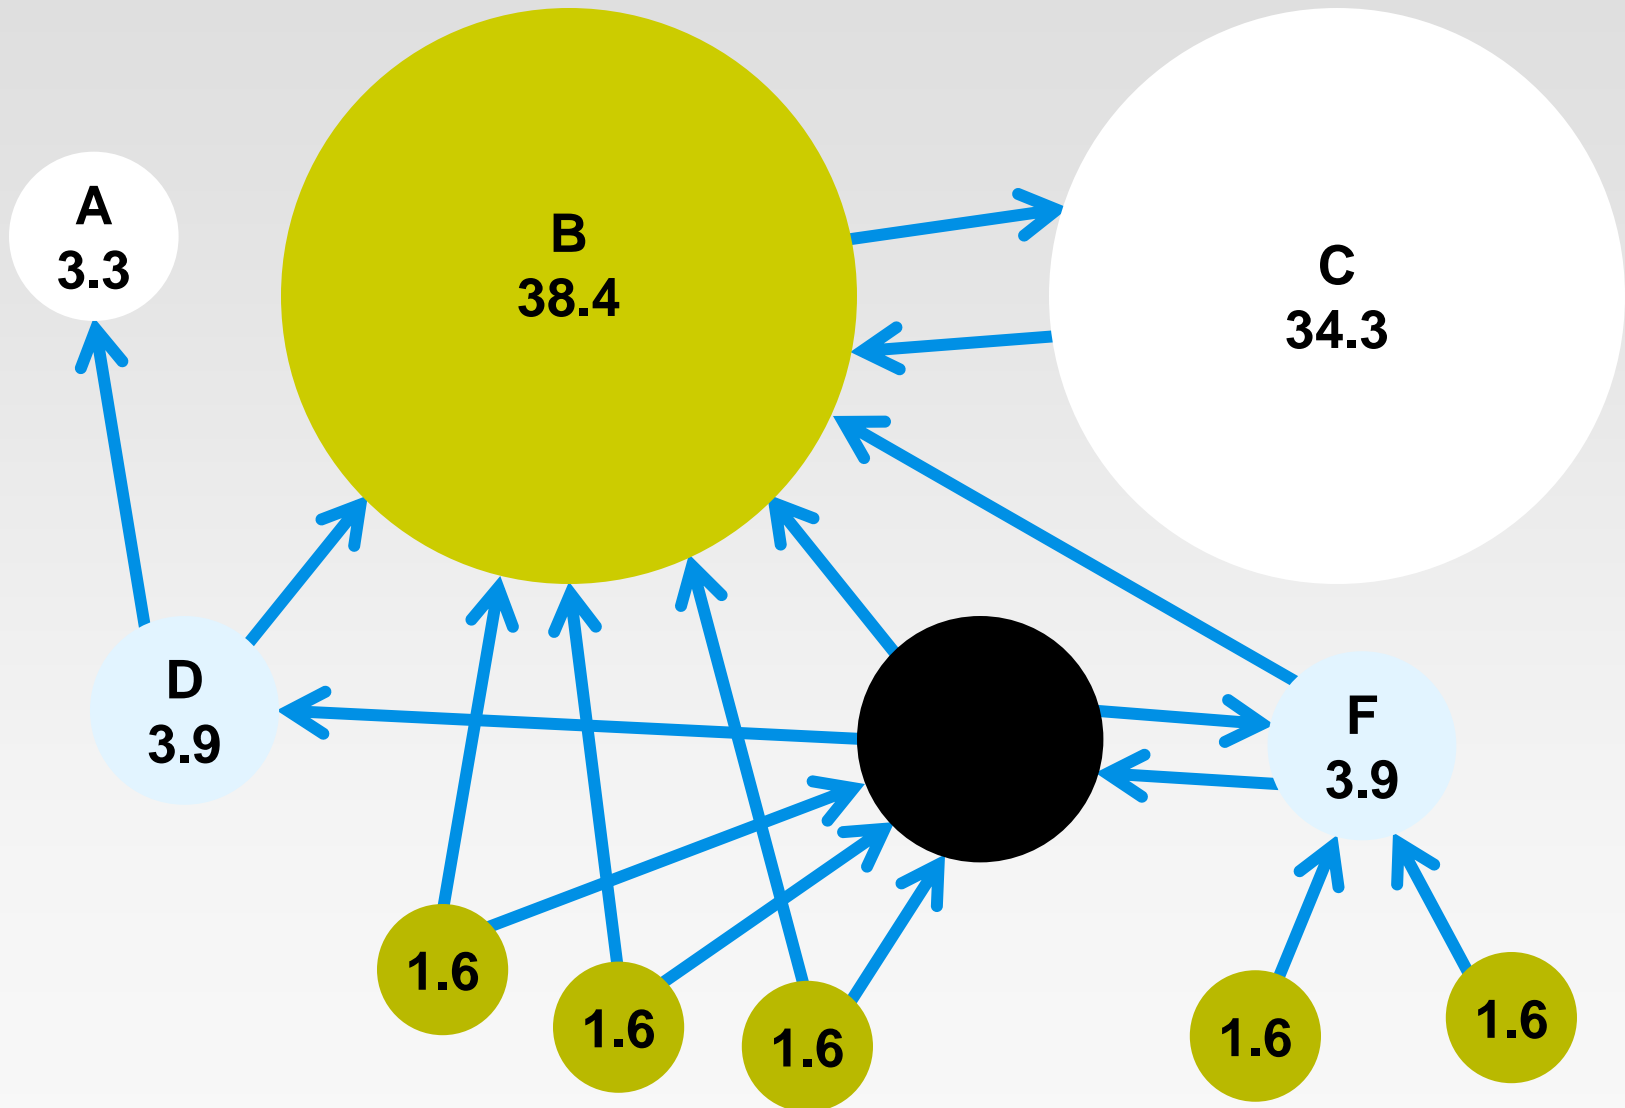
- ❖ We will cover the following Link Analysis approaches for computing importance of nodes in a graph:
  - Page Rank
  - Topic-Specific (Personalized) Page Rank
  - HITS
  - ... ..

# Links as Votes

- ❖ Idea: Links as votes
  - Page is more important if it has more links
  - In-coming links? Out-going links?
- ❖ Think of in-links as votes:
  - <http://www.unsw.edu.au/> has 23,400 in-links
  - <http://xxx.github.io/> has 1 in-link
- ❖ Are all in-links equal?
  - Links from important pages count more
  - Recursive question!




# Example: PageRank Scores



# Simple Recursive Formulation

- ❖ Each link's vote is proportional to the **importance** of its source page
- ❖ If page  $j$  with importance  $r_j$  has  $n$  out-links, each link gets  $r_j/n$  votes
- ❖ Page  $j$ 's own importance is the sum of the votes on its in-links




# PageRank: The “Flow” Model

- ❖ A “vote” from an important page is worth more
- ❖ A page is important if it is pointed to by other important pages
- ❖ Define a “rank”  $r_j$  for page  $j$

$$r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$$

$d_i$  ... out-degree of node  $i$



“Flow” equations:

$$r_y = r_y/2 + r_a/2$$

$$r_a = r_y/2 + r_m$$

$$r_m = r_a/2$$

# Solving the Flow Equations

Flow equations:

$$r_y = r_y/2 + r_a/2$$

$$r_a = r_y/2 + r_m$$

$$r_m = r_a/2$$

- ❖ 3 equations, 3 unknowns, no constants

- No unique solution
- All solutions equivalent modulo the scale factor

- ❖ Additional constraint forces uniqueness:

- $r_y + r_a + r_m = 1$

- **Solution:**  $r_y = \frac{2}{5}, r_a = \frac{2}{5}, r_m = \frac{1}{5}$

- ❖ Gaussian elimination method works for small examples, but we need a better method for large web-size graphs
- ❖ We need a new formulation!

# PageRank: Matrix Formulation

- ❖ Stochastic adjacency matrix  $M$

- Let page  $i$  has  $d_i$  out-links

- If  $i \rightarrow j$ , then  $M_{ji} = \frac{1}{d_i}$  else  $M_{ji} = 0$

- ▶  $M$  is a **column stochastic matrix**

- Columns sum to 1

- ❖ **Rank vector  $r$ :** vector with an entry per page

- $r_i$  is the importance score of page  $i$

- $\sum_i r_i = 1$

- ❖ The flow equations can be written

$$r = M \cdot r$$


# Example

❖ Remember the flow equation:  $r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$

❖ Flow equation in the matrix form

$$M \cdot r = r$$

➤ Suppose page  $i$  links to 3 pages, including  $j$



# Eigenvector Formulation

- ❖ The flow equations can be written

$$\mathbf{r} = \mathbf{M} \cdot \mathbf{r}$$


- ❖ So the **rank vector**  $\mathbf{r}$  is an **eigenvector** of the stochastic web matrix  $\mathbf{M}$

- In fact, its first or principal eigenvector, with corresponding eigenvalue  $1$ 
  - ▶ Largest eigenvalue of  $\mathbf{M}$  is  $1$  since  $\mathbf{M}$  is column stochastic (with non-negative entries)
    - *We know  $\mathbf{r}$  is unit length and each column of  $\mathbf{M}$  sums to one, so  $\mathbf{M}\mathbf{r} \leq \mathbf{1}$*

**NOTE:**  $\mathbf{x}$  is an eigenvector with the corresponding eigenvalue  $\lambda$  if:  
 $\mathbf{A}\mathbf{x} = \lambda\mathbf{x}$

- ❖ We can now efficiently solve for  $\mathbf{r}$ !
  - The method is called Power iteration

# Example: Flow Equations & M



	y	a	m
y	1/2	1/2	0
a	1/2	0	1
m	0	1/2	0

$$r_y = r_y/2 + r_a/2$$

$$r_a = r_y/2 + r_m$$

$$r_m = r_a/2$$

$$r = M \cdot r$$

$$\begin{bmatrix} y \\ a \\ m \end{bmatrix} = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix} \begin{bmatrix} y \\ a \\ m \end{bmatrix}$$



# Power Iteration Method

- ❖ Given a web graph with  $n$  nodes, where the nodes are pages and edges are hyperlinks

- ❖ **Power iteration:** a simple iterative scheme

- Suppose there are  $N$  web pages

- Initialize:  $\mathbf{r}^{(0)} = [1/N, \dots, 1/N]^T$

- Iterate:  $\mathbf{r}^{(t+1)} = \mathbf{M} \cdot \mathbf{r}^{(t)}$

- Stop when  $\|\mathbf{r}^{(t+1)} - \mathbf{r}^{(t)}\|_1 < \varepsilon$

$$r_j^{(t+1)} = \sum_{i \rightarrow j} \frac{r_i^{(t)}}{d_i}$$

$d_i$  .... out-degree of node  $i$


$\|\mathbf{x}\|_1 = \sum_{1 \leq i \leq N} |x_i|$  is the  $\mathbf{L}_1$  norm

Can use any other vector norm, e.g., Euclidean

# PageRank: How to solve?

## ❖ Power Iteration:

- Set  $r_j = 1/N$
- 1:  $r'_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$
- 2:  $r = r'$
- Goto 1



	y	a	m
y	1/2	1/2	0
a	1/2	0	1
m	0	1/2	0

$$r_y = r_y/2 + r_a/2$$

$$r_a = r_y/2 + r_m$$

$$r_m = r_a/2$$

## ❖ Example:

$$\begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix} = \begin{bmatrix} 1/3 & & & & 6/15 \\ 1/3 & & & & 6/15 \\ 1/3 & & & & 3/15 \end{bmatrix} \dots$$

Iteration 0, 1, 2, ...


# Random Walk Interpretation

## ❖ Imagine a random web surfer:

- At any time  $t$ , surfer is on some page  $i$
- At time  $t + 1$ , the surfer follows an out-link from  $i$  uniformly at random
- Ends up on some page  $j$  linked from  $i$
- Process repeats indefinitely

## ❖ Let:

- $\mathbf{p}(t)$  ... vector whose  $i^{\text{th}}$  coordinate is the prob. that the surfer is at page  $i$  at time  $t$
- So,  $\mathbf{p}(t)$  is a probability distribution over pages



# The Stationary Distribution

## ❖ Where is the surfer at time $t+1$ ?

- Follows a link uniformly at random


$$p(t+1) = M \cdot p(t)$$

- ❖ Suppose the random walk reaches a state  $p(t+1) = M \cdot p(t) = p(t)$

then  $p(t)$  is **stationary distribution** of a random walk

- ❖ **Our original rank vector**  $r$  satisfies  $r = M \cdot r$

- **So,  $r$  is a stationary distribution for the random walk**



# Existence and Uniqueness

- ❖ A central result from the theory of random walks (a.k.a. Markov processes):


For graphs that satisfy **certain conditions**, the **stationary distribution is unique** and eventually will be reached no matter what the initial probability distribution at time  $t = 0$

# PageRank: Two Questions

$$r_j^{(t+1)} = \sum_{i \rightarrow j} \frac{r_i^{(t)}}{d_i} \quad \text{or equivalently} \quad r = Mr$$

- ❖ Does this converge?
- ❖ Does it converge to what we want?

# Does this converge?




$$r_j^{(t+1)} = \sum_{i \rightarrow j} \frac{r_i^{(t)}}{d_i}$$

❖ Example:

$r_a$	1	0	1	0	...
$r_b$	0	1	0	1	...

Iteration 0, 1, 2, ...

# Does it converge to what we want?



$$r_j^{(t+1)} = \sum_{i \rightarrow j} \frac{r_i^{(t)}}{d_i}$$

❖ Example:

$r_a$	1	0	0	0
$r_b$	0	1	0	0


Iteration 0, 1, 2, ...



# PageRank: Problems

## 2 problems:


- ❖ (1) Some pages are **dead ends** (have no out-links)
  - Random walk has “nowhere” to go to
  - Such pages cause importance to “leak out”
- ❖ (2) **Spider traps:** (all out-links are within the group)
  - Random walked gets “stuck” in a trap
  - And eventually spider traps absorb all importance



# Problem: Spider Traps

## ❖ Power Iteration:

- Set  $r_j = 1$
- $r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$
- ▶ And iterate



m is a spider trap

	y	a	m
y	1/2	1/2	0
a	1/2	0	0
m	0	1/2	1

$$\mathbf{r}_y = \mathbf{r}_y / 2 + \mathbf{r}_a / 2$$

$$\mathbf{r}_a = \mathbf{r}_y / 2$$

$$\mathbf{r}_m = \mathbf{r}_a / 2 + \mathbf{r}_m$$

## ❖ Example:


$$\begin{bmatrix} \mathbf{r}_y \\ \mathbf{r}_a \\ \mathbf{r}_m \end{bmatrix} = \begin{array}{ccccc} 1/3 & 2/6 & 3/12 & 5/24 & 0 \\ 1/3 & 1/6 & 2/12 & 3/24 & \dots & 0 \\ 1/3 & 3/6 & 7/12 & 16/24 & & 1 \end{array}$$

Iteration 0, 1, 2, ...

All the PageRank score gets “trapped” in node m.

# Solution: Teleport!


- ❖ The Google solution for spider traps: **At each time step, the random surfer has two options**
  - With prob.  $\beta$ , follow a link at random
  - With prob.  $1-\beta$ , jump to some random page
  - Common values for  $\beta$  are in the range 0.8 to 0.9
- ❖ Surfer will teleport out of spider trap within a few time steps



# Problem: Dead Ends

## ❖ Power Iteration:

- Set  $r_j = 1$
- $r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$
- ▶ And iterate



	y	a	m
y	1/2	1/2	0
a	1/2	0	0
m	0	1/2	0

$$\mathbf{r}_y = \mathbf{r}_y / 2 + \mathbf{r}_a / 2$$

$$\mathbf{r}_a = \mathbf{r}_y / 2$$

$$\mathbf{r}_m = \mathbf{r}_a / 2$$

## ❖ Example:



$$\begin{bmatrix} \mathbf{r}_y \\ \mathbf{r}_a \\ \mathbf{r}_m \end{bmatrix} = \begin{array}{ccccccc} 1/3 & 2/6 & 3/12 & 5/24 & & 0 \\ 1/3 & 1/6 & 2/12 & 3/24 & \dots & 0 \\ 1/3 & 1/6 & 1/12 & 2/24 & & 0 \end{array}$$

Iteration 0, 1, 2, ...

Here the PageRank “leaks” out since the matrix is not stochastic.

# Solution: Always Teleport!

- ❖ Teleports: Follow random teleport links with probability 1.0 from dead-ends
  - Adjust matrix accordingly



# Why Teleports Solve the Problem?

Why are dead-ends and spider traps a problem and why do teleports solve the problem?

- ❖ **Spider-traps** are not a problem, but with traps PageRank scores are **not** what we want
  - **Solution:** Never get stuck in a spider trap by teleporting out of it in a finite number of steps
- ❖ **Dead-ends** are a problem
  - The matrix is not column stochastic so our initial assumptions are not met
  - **Solution:** Make matrix column stochastic by always teleporting when there is nowhere else to go

# Google's Solution: Random Teleports

## ❖ Google's solution that does it all:

At each step, random surfer has two options:

- With probability  $\beta$ , follow a link at random
- With probability  $1-\beta$ , jump to some random page

## ❖ PageRank equation [Brin-Page, 98]

$$r_j = \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + (1 - \beta) \frac{1}{N}$$

$d_i$  ... out-degree  
of node  $i$

This formulation assumes that  $M$  has no dead ends. We can either preprocess matrix  $M$  to remove all dead ends or explicitly follow random teleport links with probability 1.0 from dead-ends.

# The Google Matrix

- ❖ PageRank equation [Brin-Page, '98]

$$r_j = \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + (1 - \beta) \frac{1}{N}$$

- ❖ The Google Matrix  $A$ :

$$A = \beta M + (1 - \beta) \left[ \frac{1}{N} \right]_{N \times N}$$

$[1/N]_{N \times N}$ ...  $N$  by  $N$  matrix  
where all entries are  $1/N$


- ❖ We have a recursive problem:  $\mathbf{r} = A \cdot \mathbf{r}$   
And the Power method still works!

- ❖ What is  $\beta$ ?

➤ In practice  $\beta = 0.8, 0.9$  (make 5 steps on avg., jump)



# Random Teleports ( $\beta = 0.8$ )



$$\begin{matrix} & \mathbf{M} & & \mathbf{[1/N]_{N \times N}} \\ 0.8 & \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix} & + 0.2 & \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix} \\ & & & \mathbf{A} \\ & \begin{matrix} y & 7/15 & 7/15 & 1/15 \\ a & 7/15 & 1/15 & 1/15 \\ m & 1/15 & 7/15 & 13/15 \end{matrix} & & \end{matrix}$$

y		1/3	0.33	0.24	0.26		7/33
a	=	1/3	0.20	0.20	0.18	...	5/33
m		1/3	0.46	0.52	0.56		21/33

# Computing Page Rank

- ❖ Key step is matrix-vector multiplication

- $\mathbf{r}^{\text{new}} = \mathbf{A} \cdot \mathbf{r}^{\text{old}}$

- ❖ Easy if we have enough main memory to hold  $\mathbf{A}$ ,  $\mathbf{r}^{\text{old}}$ ,  $\mathbf{r}^{\text{new}}$

- ❖ Say  $N = 1$  billion pages

- We need 4 bytes for each entry (say)
  - 2 billion entries for vectors, approx 8GB
  - Matrix  $\mathbf{A}$  has  $N^2$  entries
    - ▶  $10^{18}$  is a large number!

$$\mathbf{A} = \beta \cdot \mathbf{M} + (1-\beta) [\mathbf{1}/N]_{N \times N}$$

$$\mathbf{A} = 0.8 \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix} + 0.2 \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix}$$

$$= \begin{bmatrix} 7/15 & 7/15 & 1/15 \\ 7/15 & 1/15 & 1/15 \\ 1/15 & 7/15 & 13/15 \end{bmatrix}$$

# Matrix Formulation

- ❖ Suppose there are  $N$  pages
- ❖ Consider page  $i$ , with  $d_i$  out-links
- ❖ We have  $M_{ji} = 1/d_i$  when  $i \rightarrow j$   
and  $M_{ji} = 0$  otherwise
- ❖ The random teleport is equivalent to:
  - Adding a **teleport link** from  $i$  to every other page and setting transition probability to  $(1-\beta)/N$
  - Reducing the probability of following each out-link from  $1/d_i$  to  $\beta/d_i$
  - **Equivalent:** Tax each page a fraction  $(1-\beta)$  of its score and redistribute evenly

# Rearranging the Equation

$$\diamond r = A \cdot r, \text{ where } A_{ji} = \beta M_{ji} + \frac{1-\beta}{N}$$

$$\diamond r_j = \sum_{i=1}^N A_{ji} \cdot r_i$$

$$\begin{aligned} \diamond r_j &= \sum_{i=1}^N \left[ \beta M_{ji} + \frac{1-\beta}{N} \right] \cdot r_i \\ &= \sum_{i=1}^N \beta M_{ji} \cdot r_i + \frac{1-\beta}{N} \sum_{i=1}^N r_i \end{aligned}$$

$$= \sum_{i=1}^N \beta M_{ji} \cdot r_i + \frac{1-\beta}{N} \quad \text{since } \sum r_i = 1$$

$$\diamond \text{ So we get: } r = \beta M \cdot r + \left[ \frac{1-\beta}{N} \right]_N$$

**Note:** Here we assumed **M**  
has no dead-ends

$[x]_N$  ... a vector of length  $N$  with all entries  $x$

# Sparse Matrix Formulation

- ❖ We just rearranged the **PageRank equation**

$$r = \beta M \cdot r + \left[ \frac{1 - \beta}{N} \right]_N$$

- ▶ where  $[(1-\beta)/N]_N$  is a vector with all  $N$  entries  $(1-\beta)/N$

- ❖  $M$  is a **sparse matrix!** (with no dead-ends)

- 10 links per node, approx  $10N$  entries

- ❖ So in each iteration, we need to:

- Compute  $r^{\text{new}} = \beta M \cdot r^{\text{old}}$

- Add a constant value  $(1-\beta)/N$  to each entry in  $r^{\text{new}}$

- ▶ **Note if  $M$  contains dead-ends then  $\sum_j r_j^{\text{new}} < 1$  and we also have to renormalize  $r^{\text{new}}$  so that it sums to 1**

# PageRank: The Complete Algorithm

## ❖ Input: Graph $G$ and parameter $\beta$

- Directed graph  $G$  (can have **spider traps** and **dead ends**)
- Parameter  $\beta$

## ❖ Output: PageRank vector $r^{new}$

- **Set:**  $r_j^{old} = \frac{1}{N}$

- **repeat until convergence:**  $\sum_j |r_j^{new} - r_j^{old}| > \varepsilon$

- ▶  $\forall j: r_j'^{new} = \sum_{i \rightarrow j} \beta \frac{r_i^{old}}{d_i}$

- $r_j'^{new} = 0$  if in-degree of  $j$  is 0

- ▶ **Now re-insert the leaked PageRank:**

- $\forall j: r_j^{new} = r_j'^{new} + \frac{1-S}{N}$       **where:**  $S = \sum_j r_j'^{new}$

- ▶  $r^{old} = r^{new}$

If the graph has no dead-ends then the amount of leaked PageRank is  $1-\beta$ . But since we have dead-ends the amount of leaked PageRank may be larger. We have to explicitly account for it by computing  $S$ .

# Sparse Matrix Encoding

- ❖ Encode sparse matrix using only nonzero entries
  - Space proportional roughly to number of links
  - Say  $10N$ , or  $4 \times 10^1$  billion = 40GB
  - **Still won't fit in memory, but will fit on disk**

source node	degree	destination nodes
0	3	1, 5, 7
1	5	17, 64, 113, 117, 245
2	2	13, 23

# Basic Algorithm: Update Step

- ❖ Assume enough RAM to fit  $r^{new}$  into memory
  - Store  $r^{old}$  and matrix **M** on disk
- ❖ 1 step of power-iteration is:

**Initialize** all entries of  $r^{new} = (1-\beta) / N$

For each page  $i$  (of out-degree  $d_i$ ):

Read into memory:  $i, d_i, dest_1, \dots, dest_{d_i}, r^{old}(i)$

For  $j = 1 \dots d_i$

$r^{new}(dest_j) += \beta r^{old}(i) / d_i$

0	
1	
2	
3	
4	
5	

$r^{new}$

source	degree	destination
0	3	1, 5, 6
1	4	17, 64, 113, 117
2	2	13, 23


$r^{old}$



# Analysis

- ❖ Assume enough RAM to fit  $r^{new}$  into memory
  - Store  $r^{old}$  and matrix  $M$  on disk
- ❖ In each iteration, we have to:
  - Read  $r^{old}$  and  $M$
  - Write  $r^{new}$  back to disk
  - **Cost per iteration of Power method:**  
 **$= 2|r| + |M|$**
- ❖ **Question:**
  - What if we could not even fit  $r^{new}$  in memory?
  - Split  $r^{new}$  into blocks. Details ignored

# Some Problems with Page Rank

- ❖ **Measures generic popularity of a page**
  - Biased against topic-specific authorities
  - **Solution:** Topic-Specific (Personalized) PageRank (**next**)
- ❖ **Uses a single measure of importance**
  - Other models of importance
  - **Solution:** Hubs-and-Authorities

# PageRank in MapReduce

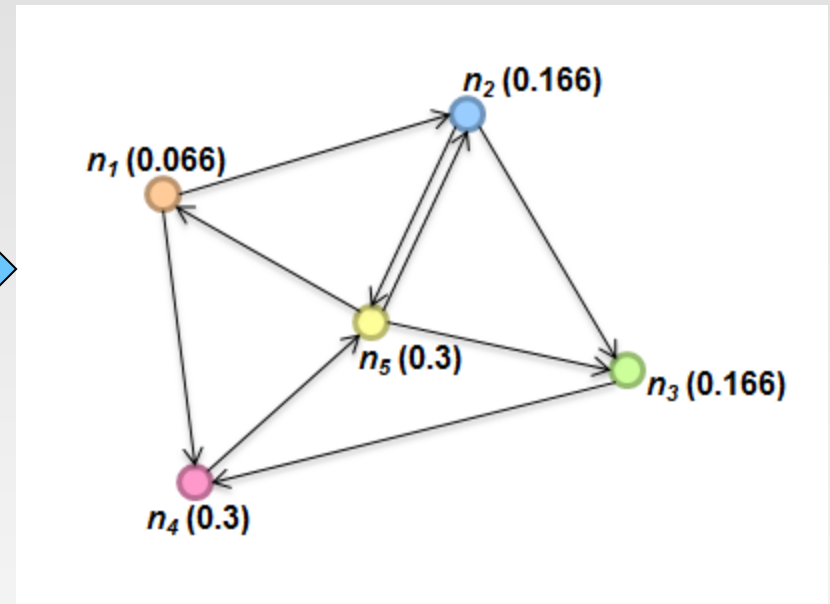


# PageRank Computation Review

- ❖ Properties of PageRank
  - Can be computed iteratively
  - Effects at each iteration are local
- ❖ Sketch of algorithm:
  - Start with seed  $r_i$  values
  - Each page distributes  $r_i$  “credit” to all pages it links to
  - Each target page  $t_j$  adds up “credit” from multiple in-bound links to compute  $r_j$
  - Iterate until values converge



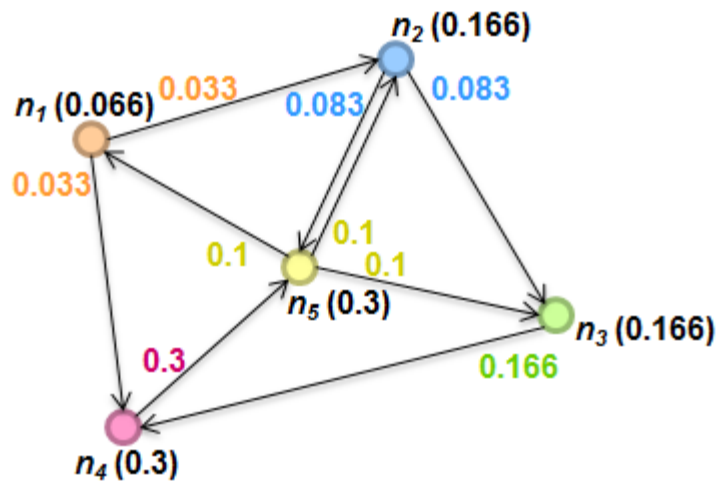
# Simplified PageRank

- ❖ First, tackle the simple case:
  - No teleport
  - No dead ends
- ❖ Then, factor in these complexities...
  - How to deal with the teleport probability?
  - How to deal with dead ends?

# Sample PageRank Iteration (1)



# Sample PageRank Iteration (2)



# PageRank in MapReduce

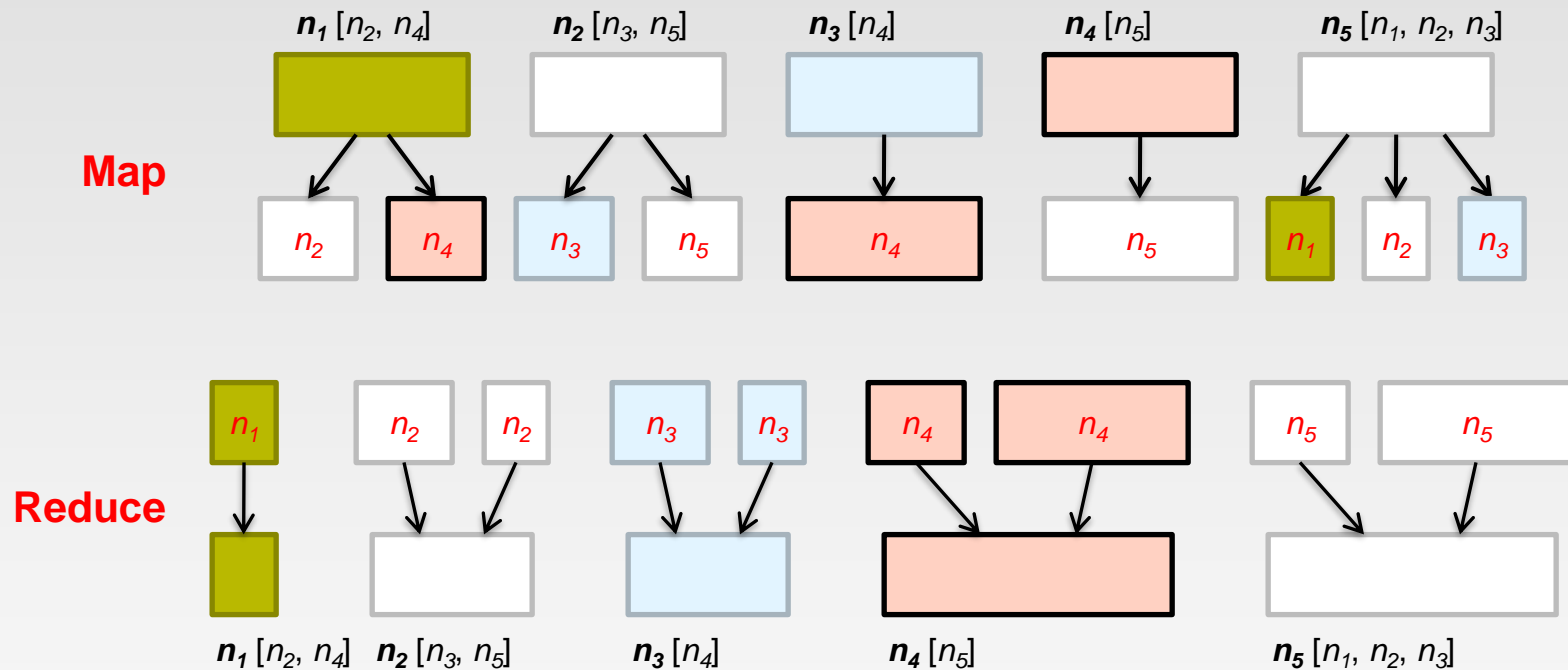
- ❖ One iteration of the PageRank algorithm involves taking an estimated PageRank vector  $r$  and computing the next estimate  $r'$  by

$$r = \beta M \cdot r + \left[ \frac{1 - \beta}{N} \right]_N$$

- ❖ Mapper: input – a line containing node  $u$ ,  $r_u$ , a list of out-going neighbors of  $u$ 
  - For each neighbor  $v$ , emit( $v$ ,  $r_u/\text{deg}(u)$ )
  - Emit ( $u$ , a list of out-going neighbors of  $u$ )
- ❖ Reducer: input – (node  $v$ , a list of values  $\langle r_u/\text{deg}(u), \dots \rangle$ )
  - Aggregate the results according to the equation to compute  $r'_v$
  - Emit node  $v$ ,  $r'_v$ , a list of out-going neighbors of  $v$



# PageRank in MapReduce (One Iteration)



# PageRank Pseudo-Code

```
1: class MAPPER
2:   method MAP(nid  $n$ , node  $N$ )
3:      $p \leftarrow N.\text{PAGERANK} / |N.\text{ADJACENCYLIST}|$ 
4:      $\text{EMIT}(\text{nid } n, N)$  ▷ Pass along graph structure
5:     for all nodeid  $m \in N.\text{ADJACENCYLIST}$  do
6:        $\text{EMIT}(\text{nid } m, p)$  ▷ Pass PageRank mass to neighbors
7:
8: class REDUCER
9:   method REDUCE(nid  $m$ , [ $p_1, p_2, \dots$ ])
10:     $M \leftarrow \emptyset$ 
11:    for all  $p \in \text{counts } [p_1, p_2, \dots]$  do
12:      if  $\text{ISNODE}(p)$  then
13:         $M \leftarrow p$  ▷ Recover graph structure
14:      else
15:         $s \leftarrow s + p$  ▷ Sums incoming PageRank contributions
16:     $M.\text{PAGERANK} \leftarrow s$ 
17:     $\text{EMIT}(\text{nid } m, \text{node } M)$ 
```

# Complete PageRank

- ❖ Two additional complexities
  - What is the proper treatment of dangling nodes?
  - How do we factor in the random jump factor?
- ❖ Solution:
  - If a node's adjacency list is empty, distribute its value to all nodes evenly.
    - ▶ In mapper, for such a node  $i$ , emit  $(\text{nid } m, r_i/N)$  for each node  $m$  in the graph
  - Add the teleport value
    - ▶ In reducer,  $M.\text{PageRank} = \beta * s + (1 - \beta) / N$

# Graphs and MapReduce


- ❖ Graph algorithms typically involve:
  - Performing computations at each node: based on node features, edge features, and local link structure
  - Propagating computations: “traversing” the graph
- ❖ Generic recipe:
  - Represent graphs as adjacency lists
  - Perform local computations in mapper
  - Pass along partial results via outlinks, keyed by destination node
  - Perform aggregation in reducer on inlinks to a node
  - Iterate until convergence: controlled by external “driver”
  - Don’t forget to pass the graph structure between iterations

# Issues with MapReduce on Graph Processing

- ❖ MapReduce Does not support iterative graph computations:
  - External driver. Huge I/O incurs
  - No mechanism to support global data structures that can be accessed and updated by all mappers and reducers
    - ▶ Passing information is only possible within the local graph structure – through adjacency list
    - ▶ Some “wasted” computations are needed
- ❖ MapReduce algorithms are often impractical on large, dense graphs.
  - The amount of intermediate data generated is on the order of the number of edges.
  - For dense graphs, MapReduce running time would be dominated by copying intermediate data across the network.


# Iterative MapReduce

- ❖ Only a subset of data needs computation:



# Iterative MapReduce

- ❖ System is not optimized for iteration:



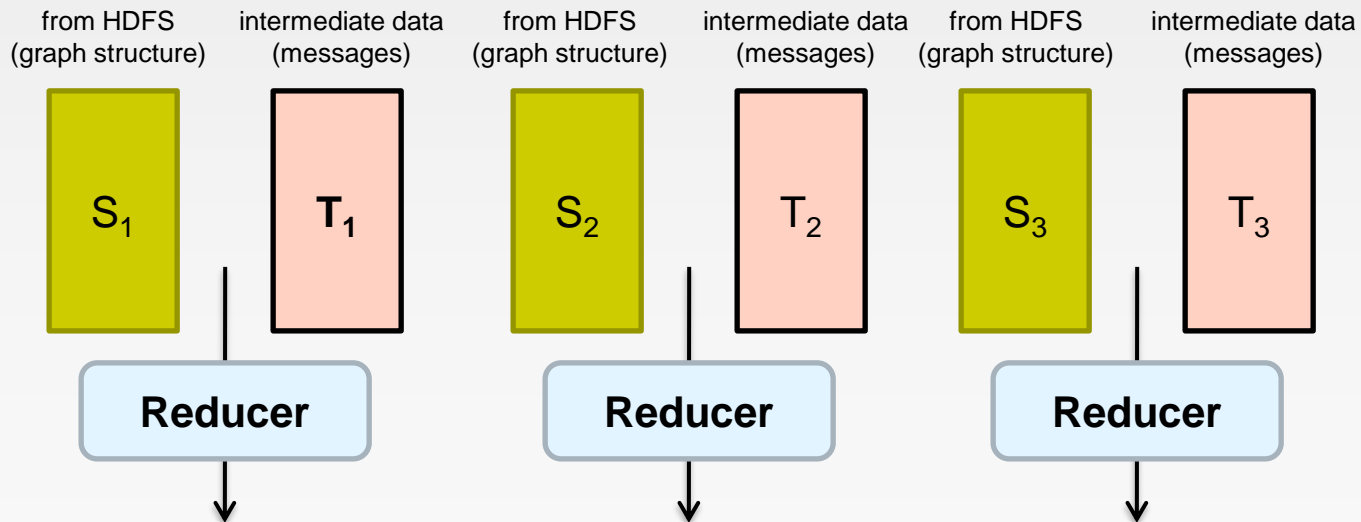
# Better Partitioning

- ❖ Default: hash partitioning
  - Randomly assign nodes to partitions
- ❖ Observation: many graphs exhibit local structure
  - E.g., communities in social networks
  - Better partitioning creates more opportunities for local aggregation
- ❖ Unfortunately, partitioning is **hard**!
  - Sometimes, chick-and-egg...
  - But cheap heuristics sometimes available
  - For webgraphs: range partition on domain-sorted URLs



# Schimmy Design Pattern

- ❖ Basic implementation contains two dataflows:
  - Messages (actual computations)
  - Graph structure (“bookkeeping”)
- ❖ Schimmy = reduce side parallel merge join between graph structure and messages
  - Consistent partitioning between input and intermediate data
  - Mappers emit only messages (actual computation)
  - Reducers read graph structure directly from HDFS



# References

- ❖ Chapter 5. Mining of Massive Datasets.
- ❖ Chapter 5. Data-Intensive Text Processing with MapReduce

**End of Chapter 9**